

CSE 522 - Real-time Embedded Systems - Spring 2021
Report - Thread event tracing in Zephyr RTOS
Shyam Joshi - 1218594676
sjoshi46@asu.edu

Implementation

1

Implementation

In thread_events.c, the following logic is used to record thread events:

```
k = irq_lock();                /* Disable interrupts for the
processor that is running the thread */
check = trace;                 /* Copy the tracing flag to a local
variable */
irq_unlock(k);                /* Enable back the interrupts for the
processor */

if(check)
{
if(tEvents_index > fix_size) {return;}
    k_tid_t id = k_current_get();
    const char *name = k_thread_name_get(id);
    threadEvents[tEvents_index].event='i';
    strcpy(threadEvents[tEvents_index].thread_name,name);
    threadEvents[tEvents_index].TSC = (temp/400);
    tEvents_index++;
}
```

1. Interrupts are disabled for the local core which is executing the thread. They are disabled because they provide atomic access. Mutex cannot be used because, whenever mutex_lock/mutex_unlock is invoked, sys_trace_void() is called and if we use mutex inside them, then it will create an infinite recursion.
2. trace flag is copied into a local variable.
3. Interrupts are enabled back.
4. tEvent_index is the variable for accessing the custom structure array which holds the trace data. It is checked if it has reached till the maximum size allowed, if yes then return.
5. Thread name is acquired using thread id of the executing thread.
6. Using tEvent_index, time stamp, thread name and event type are filled in the custom structure and tEvent_index is incremented to access the next empty custom structure array.

Custom structure is defined as follows:

```
struct threadInfo              // custom structure for recording time stamps
{
    char event;                // name of the event {i:
task_switched_in, o: task_switched_out, e: other tasks}
    char thread_name[26];      // name of the thread
    unsigned int TSC;          // time stamp variable
};
```

There are three different symbols used to indicate the type of the event and they are:

1. 'i': For the event when `sys_trace_switched_in()` is invoked.
2. 'o': For the event when `sys_trace_switched_out()` is invoked.
3. 'e': For the events such as `sys_trace_thread_create()`, `sys_trace_thread_ready()`, `sys_trace_thread_pending()`, `mutex_lock()`, `mutex_unlock()`.

`mutex_lock()` and `mutex_unlock()` is implemented in `sys_trace_void()` with the same above logic. The only difference is that a condition is placed to validate the argument with the id of event `mutex_lock` and `mutex_unlock`.

trace flag is set and unset in `sys_trace_void()` and `sys_trace_end_call()` by validating the argument of the function. They are called inside `tracing_start()` and `tracing_end()` respectively, defined in `main.c` by passing appropriate argument to start and end tracing.

In `main.c`, the following tasks are performed:

1. `NUM_MUTEXES` number of mutexes are defined.
2. `tracing_start` is invoked to start recording trace data.
3. `NUM_THREADS` number of threads are created and thread's structure is passed as an argument for its callback function.
4. Each of the threads are assigned the name defined in `task_model.h`

Same callback function is assigned to each of the threads and the following task is performed:

1. Number of iterations to be performed in each computation are stored locally.
2. A time stamp is collected before starting a busy loop.
3. A condition is checked to see if `TOTAL_TIME` amount of time is expired or not. If yes, then the thread returns.
4. Else, it performs `compute_1`, acquires lock, performs `compute_2`, releases lock and performs `compute_3`.
5. A new time stamp is collected and in the condition before entering the busy loop, it is subtracted from the time stamp collected before entering the busy loop and the result is compared with `TOTAL_TIME`.
6. When the last thread exits, `tracing_end()` and `tracing_dump()` are invoked.

The following are the two screenshots from GTKWave with `CONFIG_PRIORITY_CEILING=0` and `CONFIG_PRIORITY_CEILING=10` with the same input task set in the two runs as defined in `task_model.h`




Figure 1. GTKWave screenshot with CONFIG_PRIORITY_CEILING=0

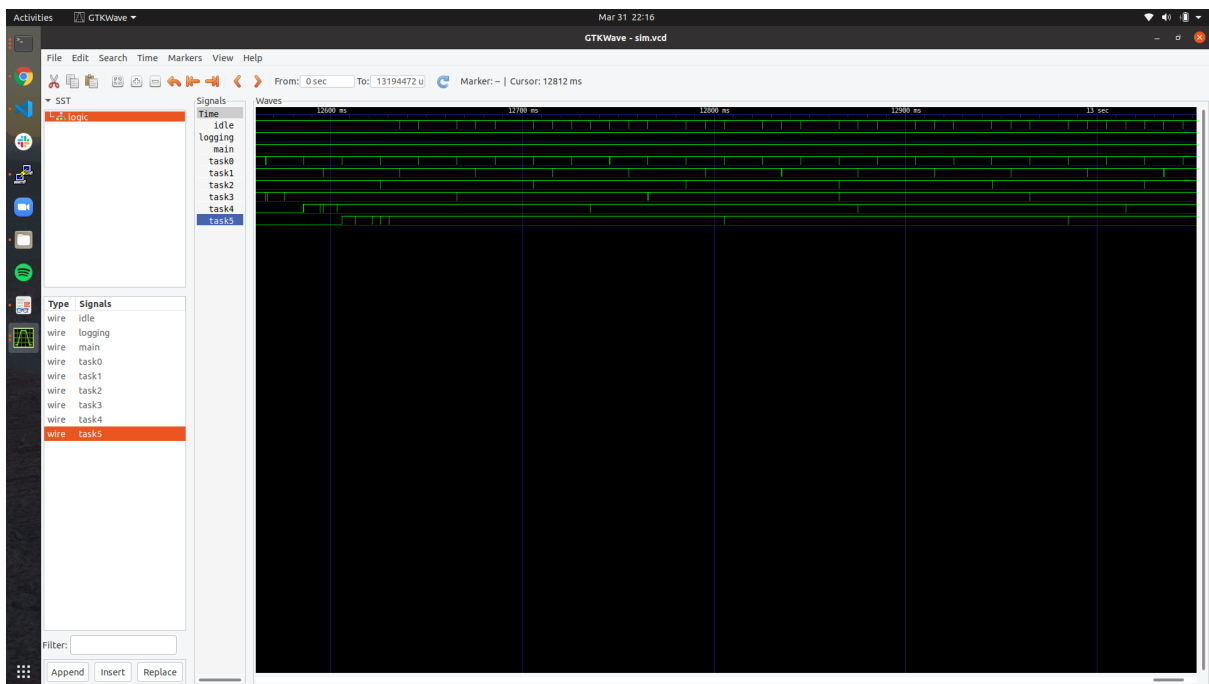


Figure 2. GTKWave screenshot with CONFIG_PRIORITY_CEILING=10